

Design Review — Feedback, Sort, and What Changed

Site reviewed: <https://yuguda.netlify.app/> Reviewer: Fuhad, backend developer Proof statement given to the reviewer (Week 1):

"I am a backend engineer who architects and builds the robust, high-performance systems required to serve AI models in production. If you are an Engineering Manager or technical recruiter looking to scale an AI-integrated product, email me to schedule a technical screening or interview."

I asked the two questions blind — link only, no proof statement — so the ten-second read wasn't contaminated by me telling him what to look for. The proof statement came after he'd answered.

The two questions

"In ten seconds, what do I do?" He got it: backend engineer moved into AI/agent pipelines — Python, FastAPI, some decentralized work, deploying open-weight models. He said the headline and subhead do their job.

"Would you believe I'm good at it?"

"Cautiously yes, but with an asterisk... in ten seconds I haven't seen a single number, outcome, or proof point — no 'reduced X,' no user count, no scale. It sounds competent. It hasn't yet shown me competent. A skeptical recruiter gives you credit for vocabulary, not yet for results."

That's the answer that mattered. The claim landed; the proof didn't.

The rest of it, unedited in substance

- My best evidence lived on a **second domain** (github.io). A recruiter skimming for 30–45 seconds does not click through, so the case studies effectively didn't exist for the read that counts.
- **Zero numbers anywhere.** Every claim qualitative and self-reported.
- The capstone was the **most convincing thing on the page** — real mechanism, explicit safety constraints, a build log with a caught bug — and it sat *below* the weaker case-study links.
- The **stack section was filler**: bragging about Netlify's free tier and pointing a CNAME "actively lowers perceived seniority... if this counts as a bullet point, what else on this page is padding?"
- **Four competing CTAs** with equal weight and no hierarchy. The lowest-friction action for a busy EM — just emailing me — was "buried at the very bottom, inline with copyright text." Asking for a 30-minute call before showing any evidence is "premature."
- **Does the work back the claim?** "Side projects — including, ironically, the one you led with." His sharpest point: the capstone is a script that sweeps my own repos, one user, me. Structurally a demo, sitting under a headline claiming "systems that ship." Meanwhile CampusPQ — which actually shipped to real users — got one clause and a link away. **I had inverted my own evidence.**
- Cardano mentioned with no outcome invites a "crypto portfolio" pattern-match, and nothing on the page defused it.
- No visuals anywhere — "a wall of text describing systems with zero visual proof they exist."

Must-fix vs. nice-to-have

Must-fix

#	Problem	Why it's must-fix
1	No numbers or outcomes anywhere	The proof doesn't land — the exact failure the proof statement exists to prevent
2	Email buried in the footer under four equal CTAs	Directly blocks the one action the proof statement names
3	Case studies were link-outs to another domain	The primary audience never sees the evidence
4		Undermines the headline claim on the same screen

#	Problem	Why it's must-fix
	Evidence inverted — solo tool showcased, shipped product buried	
5	Netlify/DNS "stack" section	Actively lowers perceived seniority; casts doubt on stronger claims
6	No visual proof of any system	Claims unsupported at a glance

Nice-to-have

- Voice drifting from the sharp opening line into generic infra-speak.
- Deeper Cardano context beyond defusing the pattern-match.
- Per-case architecture diagrams.

Not a defect — checked, not assumed

He reported the capstone video "doesn't play inline, tells the visitor to download an .mp4," and separately that my case-study site, CV, and GitHub profile "all bounced back empty on my end (tooling issue, not a comment on your content)."

I verified before acting on it rather than either dismissing or blindly "fixing" it: the file is **H.264 High profile, yuv420p** — the most broadly supported combination there is — served with `content-type: video/mp4` and `accept-ranges: bytes`. Rendered in a real browser it loads and shows controls at `0:00 / 1:10`. His tool was rendering the `<video>` fallback text because it doesn't execute media, which is the same reason three other links came back empty for him. Same root cause, so I treated it as one tooling artifact rather than four separate bugs.

I did add `playsinline` regardless, since iOS Safari will otherwise force fullscreen playback — a real (if unrelated) mobile issue his report made me go looking for.

What I changed on the live site

1. Numbers, on the page, above the fold. CampusPQ now leads with **5,000+ questions processed, 98% success rate, live in production**. Those figures were sitting on my CV the whole time and never made it onto the site — which is the part that stings.

2. Email is now the single primary action. One large gradient button, `Email me for a technical screening`, with a prefilled subject line. Calendly is demoted to "Or book a slot," CV/GitHub/ LinkedIn are visually quieter, and the raw address sits under the buttons in monospace. A closing block at the end of the page repeats the email CTA alone, so the ask is the last thing on screen. Full-width on mobile.

3. All three case studies moved on-page, with screenshots. No more one-line teaser plus a link to another domain. Each has the problem, the mechanism, a real screenshot (tap for full size), and its stack. The github.io write-ups are now a "longer version" link rather than the only version.

4. Evidence order inverted back. CampusPQ (shipped, real users, real numbers) leads. The capstone moved below the shipped work and is labelled **"Internal tool · one user (me)"** — I'd rather say that outright than let an EM catch the mismatch himself. Its value is reframed as *how* it's built: a deterministic script gathers the git data and the model may only summarize what that script printed, so it cannot invent activity by construction rather than by instruction.

5. Stack section deleted. Nothing about Netlify's free tier survives.

6. Status labels on every project — `Live · real users`, `Deployed · testnet`, `Workflow live · agent designed` — so nobody has to guess what's shipped versus specced.

7. Cardano defused with mechanism. Edix now states plainly that only a certificate *hash* goes on-chain, one transaction per certificate, that no student data is published, and that it runs on preprod testnet, not mainnet.

Contrast re-verified after the rework: every text colour passes WCAG AA (5.91:1 at the lowest, most well above 10:1).

What I'm not doing yet

FL-04 and Edix still have no numbers, because I don't have honest ones for them — no latency, uptime, or usage figures I'd stand behind. Rather than manufacture metrics to fill the pattern, those two lead with mechanism and carry explicit status labels. That gap is real and stays visible until I have figures worth publishing.