

The Through-Line, Identity Kit & Image Set

Yuguda Muhammed — Machine Learning Engineer Live: <https://yuguda999.github.io/ys-dev/> · <https://yuguda.netlify.app/>

1. The one-line claim

I build models that survive production.

Chosen from ten options I generated, then cut. The runners-up and why they lost:

- "Most LLM demos die on contact with real data. I build the ones that don't." — sharper opinion, but it leads with someone else's failure instead of my work, and it's two sentences where one will do.
- "I move LLM systems from 30% right to 90%+ — with a harness that proves it." — carries my hardest number, but welds the whole claim to one project's metric. If that number ages, the claim ages with it.

"Survive production" won because it names the actual dividing line in ML work. Anyone can get a model working in a notebook. The claim is that mine keep working when real users, messy inputs and concurrency arrive — and every case study on the site is evidence for exactly that sentence.

The one action it ladders to: email me for a technical screening.

2. Content map

Five pages. Every page ends in the same action.

Home — index.html

#	Section	Purpose
1	Hero	The claim, in one line, above the fold
2	Capability marquee	Scannable proof of range
3	Selected work	6 projects, strongest first
4	What I do	Three capability groups: LLM & agent systems · Training & evaluation · ML infrastructure
5	Experience	Lead AI Engineer, and prior roles
6	More projects	10 GitHub repos + Project Edix
7	Contact	CTA: Email me

Case order in Selected work — deliberate, strongest first:

- Proposal Generation Engine — in production, client work, and the only case with a before/after quality number (30% → 90%+ under an evaluation harness). Leads because it is the most senior signal on the site.
- FL-04 Orchestrator — open-weight model served locally; shows I run inference myself rather than renting an API.
- CampusPQ — shipped, real users, 5k+ questions at 98%.
- DevMemory — open source, installable from PyPI by anyone.
- Eeva — multimodal, open-weight vision model.
- Receipt Classifier — classical ML, shows the range isn't only LLMs.

Case pages — proposal · fl04 · campuspq · edix

All four follow one repeated shape, which is the point:

#	Section
1	Back link → Selected work
2	Title + one-line lede
3	Facts strip (type · scale · status · stack)
4	The problem
5	Approach / The design
6	Evaluation or Current state — <i>what came of it</i>
7	Stack
8	Prev / next case
9	Footer — CTA: Email me

Project Edix is deliberately reachable from *More projects* rather than Selected work, and says on its own page that it is systems architecture, not ML. It stays because the reasoning transfers; it does not get promoted because it would dilute the claim.

Still need to gather

Honest list, so build week isn't blocked by discovering these late:

Gap	Which page	Why it matters	Blocker?
No screenshot for Proposal Engine — my lead case	proposal.html	The strongest case is the only one with no visual proof	Client work; may not be shareable. Need to confirm what I can show — a redacted output, the eval harness table, or an architecture diagram I draw myself
No screenshot for DevMemory	index card	It's public and installable — a terminal capture of memory written in one AI client and recalled in another would prove the core claim in one image	No — I can capture this today
No photo of me anywhere	Home / contact	Brief says anything that is <i>me</i> should be a real photo, not generated	No — need to take one
Eeva + Receipt Classifier have no captures	index cards	Currently text-only cards	No
No numbers for FL-04 or Edix	those case pages	Both lead on mechanism instead. Honest, but thinner than CampusPQ and Proposal	Needs real measurement, not invention
Open Graph share image	all pages	Meta tags for title/description exist; no og:image yet, so links unfurl bare	No

3. Identity kit

Type — two fonts, both free on Google Fonts

Role	Font	Usage
Headings and body	Inter	400 / 500 / 600 / 700. Headings -0.032em tracking, line-height 1.05 ; hero pushes to -0.05em . Body 16px / 1.6
Labels, metadata, UI	JetBrains Mono	400 / 500 only. Section eyebrows, nav, tags, facts strips. 0.72rem , 0.14em tracking, uppercase

One typeface does the reading. The mono is not decoration — it marks machine-ish information (labels, counts, statuses) so structure is legible without adding a third voice.

Palette — hex codes

Dark (default)

Token	Hex	Role
--bg	#08090A	Near-black background
--card	#0C0E10	Raised surface
--ink	#F2F4EF	Near-white primary text
--ink-2	#B6BAB2	Secondary text, hero headline
--muted	#82877D	Mono labels
--accent	#BEF264	The one accent — links, focus rings, brand mark
--line	#1D2124	Hairline borders

Light (opt-in via toggle)

Token	Hex
--bg	#FAFAF7
--ink	#0A0B0C
--ink-2	#3C4046
--muted	#666B63
--accent	#4D7C0F

Four colours doing the work: near-black, near-white, one grey, one accent. The accent is the *only* saturated colour in the entire system. **Contrast, checked not assumed** (WCAG AA needs 4.5:1 for body text):

Pair	Ratio	
#F2F4EF on #08090A	17.99	PASS
#B6BAB2 on #08090A	10.11	PASS
#BEF264 on #08090A	15.25	PASS
#0A0B0C on #BEF264 (accent buttons)	15.08	PASS
#82877D on #0C0E10	5.26	PASS
#666B63 on #FAFAF7 (light)	5.22	PASS

The muted grey **failed before this week** — #74796F measured 4.47 on the background and 4.33 on cards, and the light theme's #7B8078 was 3.86. All three are under the 4.5 minimum, and that colour is used for every section label on the site. Corrected to #82877D / #666B63, which clear 5.2 with real headroom rather than scraping past at 4.51 — the brief asks specifically about reading a phone in sunlight, and 4.5 exactly is not a margin.

Logo / favicon

A lime chip with lowercase `ym` in JetBrains Mono, #0A0B0C on #BEF264, 17/64 corner radius. It is the same mark already sitting in the site header and footer, so the browser tab and the page agree.

Files: `favicon.svg` (scalable), `favicon-32.png`, `favicon-16.png`, `apple-touch-icon.png` (180), `icon-512.png`, `site.webmanifest`. Declared on all five pages.

Kept to two letters and one colour on purpose — at 16px anything more becomes mush, and the recognisable thing at that size is the lime block, not the letterforms.

Style note

Inter for everything you read, JetBrains Mono for everything you scan. Near-black #08090A, near-white #F2F4EF, one lime accent #BEF264, and nothing else saturated.

The mood is *instrument panel* — quiet, dense, technical, with generous space and a single bright colour used sparingly enough that it still means something. The lime appears on exactly one thing per screen. The work is the only place colour is allowed to crowd.

4. Image set, and what I rejected

What's on the site

Image	Page	Type	Why it earns its place
fl04-vllm-logs.png	FL-04	Real capture	vLLM serving real completions — proves I run inference locally, which is the whole claim of that case
campusq-api-docs.png	CampusPQ	Real capture	Live OpenAPI surface — proves the service exists and is deployed
edix-blockchain-identity.png	Edix	Real capture	Wallet on Cardano preprod, honestly labelled as testnet
ym favicon	all	Made	Two letters, one colour
Typographic posters	Proposal, DevMemory cards	Made — CSS, not an image file	See below

All three work captures are real screenshots, downscaled to ≤1400px and quantized (741 KB → 348 KB across the set), and each is wrapped in a link so a phone user can open it full size instead of squinting at a shrunken 1400px capture.

The rejection note

I rejected the entire category of AI-generated hero imagery, and never generated a single one.

The pull was real — the Proposal Engine is my lead case and it has no screenshot, so there's an obvious empty slot begging for an abstract "neural network" render. I didn't fill it, for two reasons.

First, generated abstract art for an ML portfolio is the exact thing a reviewer has seen a hundred times: the melted glass, the glowing blue brain, the impossible geometry. It reads as decoration hiding an absence, and a technical reader clocks it instantly — *if there were a real screenshot, it would be here*.

Second, and worse: putting a fake picture of a system next to claims about that system is a credibility trade I'd lose. The whole positioning is "I can prove it works." A generated image is unprovable by construction.

What went in the empty slots instead: a flat typographic poster — the project's one-sentence claim in the site's own type, on a single dark tint, built in CSS with no image file at all. It's honest about being a label rather than evidence, it costs zero bytes, and it stays consistent because it's made of the same two fonts and one accent as everything else. An empty slot handled with type looks deliberate; an empty slot handled with AI slop looks like it's hiding something.

One judgment call still open

The hero runs an animated network-graph canvas in the accent colour. It's procedural, not generated — drawn in `bg.js` — and it disables itself under `prefers-reduced-motion`. But it is the most eye-catching thing above the fold, and this week's rule is that the frame must not steal from the painting. It survives for now because it sits behind the text at low opacity and reads as a plot rather than an ornament. If a reviewer says it competes with the work, it goes — that's a one-line change.